



Πολυτεχνείο  
Κρήτης

# **CoRAM++: Supporting Data-Structure-Specific Memory Interfaces for FPGA Computing**

Τρίμας Χρήστος  
Πίσκοπος Διονύσης

LAB59142109

# Contents

- Introduction
- Motivation
- CoRAM vs CoRAM++
- Background and Related Work
- Supported Data Structures and Interfaces
- Implementation of the Interfaces
- Evaluation
- Results

# Introduction


- Facilitating DRAM access is an essential part of an application programming environment for FPGA computing.
- Existing FPGA applications support only simple memory access patterns.
- Introducing CoRAM++, a programming environment for FPGA.

# Motivation

- Accessing data in external DRAM, essential for FPGA computing.
- Before CoRAM++, application level interfaces to simplify interaction with DRAM and enable application portability.
- In need for application level interface that supports complex data structures.

# CoRAM/CoRAM++

- Decomposition of an FPGA application.
- CoRAM: Protection of the kernels. How?
- Not direct communication between Hardware kernels and DRAM.



# CoRAM/CoRAM++

- CoRAM penalties in today's FPGA applications:
  - Reduced performance.
  - Logic resource overhead.
- How CoRAM++ handles those problems?
  - Reconfigurability.
  - More customized datapath to memory.
  - Extensible library of data structure specific data transport modules.
- Example: To stream a block of data, provision of specialized commands, initiating, monitoring data streaming between DRAM and hardware kernels. Similarly specialized for other data structures cases.

# Background and Related Work

- Altera, Xilinx similar system builder tools for system integration.
- MaxCompiler (Maxeler), decomposed applications in to kernels and manager modules, for computation and I/O, Java-like language.
- RedSharc similar to CoRAM/CoRAM++, except for global controlling by software on embedded processor.
- LEAP, scratchpad memory environment backed up by off chip-DRAM. Multiple hardware kernels, can share the same virtualized scratchpad. Disadvantage: Generic abstractions, no specific access patterns optimization.

# Supported Data Structures

- **Block Copy.**  
Random access by the hardware kernel to data within a block. (CoRAM)
- **Read/Write Stream.**  
Sequential single data element access by the hardware.
- **Multi-Dimensional Array.**  
Pre-defined size dimension, tiled lay out optimization, sequential access(Read/Write data access pattern), reordering data support for permutation engines.
- **Linked List.**  
Compile-time definition of the data structure, allows sequential access(Read/Write data access pattern).



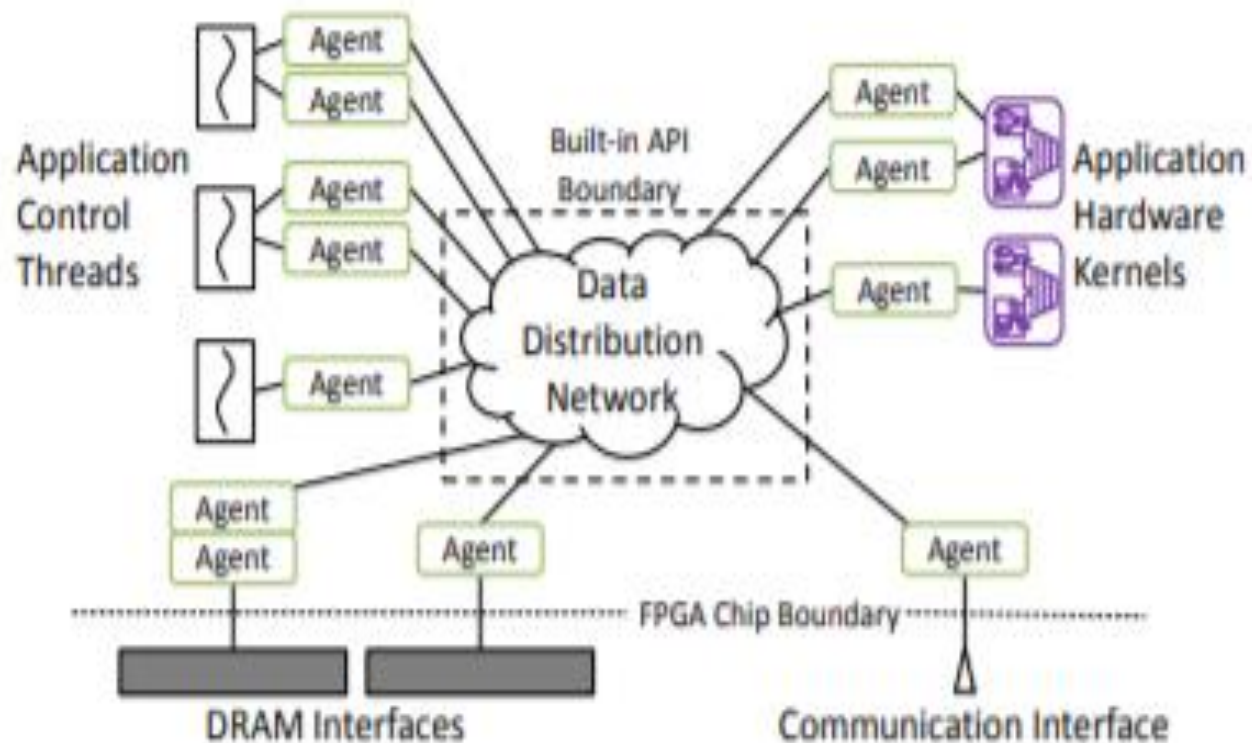
# Additional application level Interfaces

- Control thread side Scratchpad, access data from DRAM, based on data locality.
- Host Computer manages data transfer, between a host computer and FPGA's DRAM.
- Channel interface transmits messages between hardware kernels and control threads.

# Implementation of CoRAM++ Interfaces

- Agents: Hardware modules, implementing the application level interfaces.
- Connection of control threads to agents through a command interface, agent-supporting function calls, converted in to messages to hardware kernels, or transfer to another agent's control thread for sequential operations.
- Port interfaces(FIFOs, SRAMs and CAMs) for hardware kernels accordingly.
- Capability of agents to send commands to data distribution network, using fire and forget, execution model.
- Agent structure: not device specific based, instead supported by a thin hardware virtualization layer, providing the following interfaces:
  - Data Transfers.
  - Messages.
  - Direct connections for DRAM and I/O.
- Finally, combination with different programming environments such as LEAP.

# System level design of CoRAM++




# Evaluation process

- For the evaluation process, the following boards were used.
- The main purpose of the evaluation, was investigation of bandwidth bound and latency bound scenarios(not compute bound scenarios).

FPGA Board	Terasic DE4	Xilinx ML605	Xilinx ZC706
FPGA	Altera EP4SGX530	Xilinx LX240T	Xilinx XC7Z045
Hard CPU Cores	None	None	2×ARM A9
Logic Cells	531,200	241,152	350,000
Block Memory	3.4 MB	1.8 MB	2 MB
Hard DSP blocks	1,024	768	900
DRAM Bandwidth	2×6.4 GB/s	6.4 GB/s	12.8 GB/s (Fabric only) 4.3 GB/s (Shared)
FPGA Vendor Tool	Quartus 14.0	ISE 14.7	ISE 14.7

# 1<sup>st</sup> experiment.

- Spiral generated, 64-input double precision DFT engine, 200MHz, streaming input from DRAM, streaming output to the other DRAM.




Component	ALUT %	Register %	Block Mem %
DRAM Controllers	11.4	12.3	15.3
NoC	6.7	6.2	0
NoC Endpoints	16.3	23.8	25.6
Threads + Agents	9.0	12.8	44.9
Kernel	56.6	44.9	14.2

- From the above results, streaming DFT, at 512 Kbytes has 90% peak from best read/write performance(red/green line). Maximum peak at, at least 8Mbytes of data. Limited performance at small data size sets, due to the latency of the DFT core.
- Implementation Resources: 32% of the logic and 7% of the block memory, significant amount of resources.

# Multidimensional array accesses(2<sup>nd</sup> Experiment)

- Use permutation engines to recover rows from tiles of data and transpose tiles of data, in 2D and 3D DFT, each one decomposed in 1D DFT.
- Strided Array Accesses:  
1024x1024 double precision DFT, row major data. Ineffective, due to misses in the DRAM row buffer, caused by strided data accesses during column traversal.
- Tiled Array Accesses(3 configurations):  
CoRAM++ Row-inline, very slow, required the same strided accesses as the row major computation.  
Row-extrapass, better performance, extra pass through memory, after final phase of computation.  
Tiled, 90% matched the referenced calculations, keep data tiled throughout the computation. However, higher resource utilization than the reference DFT.



# Resource utilization for CoRAM++ and Reference DFTs

- While CoRAM++ has more overhead, provides flexibility and ease in the program environment, selection of array layout and array traversal method in DRAM, also hosts computer interface.

DFT Type	Logic %	Block Mem %	DSP %
CoRAM++ 1024×1024 2D	62	58	35
CoRAM++ 512×512×512 3D	52	51	20
Reference [17] 1024×1024 2D	27	47	35

# Linked List accesses(3<sup>rd</sup> Experiment)

- Linked List accesses(5 methods):

**Scratchpad-1:** C like linked list, accessing DRAM using the scratchpad agent, single row data matching the DRAM interface width.

**Scratchpad-8:** Same as scratchpad-1, 8 rows of data, same size as previous case, for locality and low cost of DRAM access.


**Scratchpad-64:** same method with 64 rows.

**Agent only:** Instantiates the hardware linked list engine within the linked list agent, linked list operations are executed from the linked list agent, triggered by a control thread(best performance for CoRAM).

**Agent+Accelerator:** Usage of the linked list agent with hardware linked list engine attached to the DRAM(Accelerator).




# Linked List Traversal



- Best case(packed linked list nodes and data into contiguous blocks) and worst case(linked list separated linked list nodes and data by 8 kilobytes) DRAM layouts.
- **Scratchpad**, improved performance by increasing the number of rows, however reduced performance in the worst case.
- **Agent only**, 1.7x faster than scratchpad for best case, slightly faster in the worst case.
- **Agent+Accelerator**, 9x faster than the baseline traversal on DE4, 8.5x on ML605 and 10.5x on ZC706, 5.2x faster than agent only.

# Sorted Linked List Merge



- 100 merged pairs of linked lists randomly assigned to each list pair, aligned to the DRAM interface width, 50% chance being a following node, or in a random location.
- Agent+Accelerator linked list configuration much faster than any other configuration, lower latency DRAM accesses when chasing pointers.

# Conclusions about the CoRAM++/Future Work

- Conveniently and efficiently support simple and complex data access patterns, through a library of data structure specific interfaces.
- Matches the performance of hand designs, using data structure specific modules connected to a DRAM interface achieves 5.2x better performance on data dependent operations.
- Future explorations will target in the acceleration of other data structures, such as trees and graphs.
- Investigation of hardware/software codesign of hybrid CPU-FPGA system and target FPGA programming environments, rather than devices

END